

Parameter fitting for damped harmonic oscillator using Stan

Konsta Parkkali, Valtteri Turkki

CS-E5710 Bayesian Data Analysis, fall 2022




Aalto University

“In this project we shall go nuts
with nuts using HMC-NUTS.”

Introduction


- Inspiration from Stan case study Lotka-Volterra*
 - Motivation: Estimating the parameters of a dynamic system using bayesian inference
- Data was hard to find → measure own data
- Studied system was simple harmonic oscillator
 - We used two different pendulums
 - Interests in estimating the damping
- Two differential equation based Bayesian models
 - Pooled and hierarchical
 - Separate model was left out due to project length



Data

$$\ddot{x} = -\frac{\mu}{ml}\dot{x} - \frac{g}{l}x$$

- We needed damped oscillations for the ODE
 - First we tried glass, but a classic pendulum was better
- Measurement setup:
 - A nut hang on the end of a string → nearly a point mass
 - Oscillation observed by taking a video using phone camera
 - The position $x(t)$ was tracked using Tracker software (<https://physlets.org/tracker/>)
 - Final data was two 700 $x(t)$ -point sets



Models - Priors

$$\ddot{x} = -\frac{\mu}{ml}\dot{x} - \frac{g}{l}x$$

- Ill-posedness is evident \rightarrow informative priors were needed
- The mass m , string length l and gravity g are well known
 - This allows regularizing the solution
- All parameters positive by construction
- Normal priors used with best estimate mean and scale given confidence
- The 90% confidence levels:
 - Mass: 12 ± 8 grams
 - String length: 17.5 ± 5 cm
 - Gravitational acceleration: 9.81 ± 0.05 m/s²
 - Initial state: $x_0 = 8 \pm 4$ mm, $\dot{x}_0 = 0 \pm 0.05$ m/s
- Air-resistance:
 - No prior information
 - ODE characteristic equation gives upper limit $\mu < \sqrt{4m^2lg}$
 - $\rightarrow \mu < 0.1$ kg*m/s

Models - Likelihood model

- Given the parameters $\phi = (m, g, \mu, l)$ and the initial condition (x_0, \dot{x}_0) for the pendulum, we can simulate the oscillation by solving the ODE
- Measured data is normally distributed from the simulated data
 - Measurement error, model approximations, imperfect oscillation
 - $x_t \sim N(\text{sim}(x_0, \dot{x}_0, \phi, t), \sigma)$, where $\sigma \sim N_+(0, 1)$
- Pooled model: pendulums have same underlying parameters, different initial conditions
- Hierarchical model: pendulums can have also different masses and air-resistances.
 - Hyperparameters μ_0 and m_0

Models - Pooled model

$x_{t,i} \sim N(\text{sim}(x_{0,i}, \dot{x}_{0,i}, \phi, t), \sigma)$ for $i = 1, 2, t = 1, \dots, 700$, and where $\phi = (m, g, \mu, l)$.

$$\sigma \sim N_+(0, 1)$$

$$m \sim N_+(0.012, 0.0049)$$

$$g \sim N_+(9.81, 0.03)$$

$$\mu \sim U(0, 0.1)$$

$$l \sim N_+(0.175, 0.03)$$

$$x_{0,i} \sim N(0.008, 0.002) \text{ for } i = 1, 2$$

$$\dot{x}_{0,i} \sim N(0.0, 0.07) \text{ for } i = 1, 2$$



Models - Hierarchical model

$x_{t,i} \sim N(\text{sim}(x_{0,i}, \dot{x}_{0,i}, \phi_i), \sigma)$ for $i = 1, 2, t = 1, \dots, 700$, and where $\phi_i = (m_i, g, \mu_i, l)$.

$$\sigma \sim N_+(0, 1)$$

$$m_0 \sim N_+(0.012, 0.0005)$$

$$m_i \sim N_+(m_0, 0.0049) \text{ for } i = 1, 2$$

$$g \sim N_+(9.81, 0.03)$$

$$\mu_0 \sim U(0, 0.1)$$

$$\mu_i \sim N_+(\mu_0, 2\mu_0) \text{ for } i = 1, 2$$

$$l \sim N_+(0.175, 0.03)$$

$$x_{0,i} \sim N(0.008, 0.002) \text{ for } i = 1, 2$$

$$\dot{x}_{0,i} \sim N(0.0, 0.07) \text{ for } i = 1, 2.$$

Implementation and running


- Models computationally heavy
 - Laptop and JupyterHub could do only 200 samples in 15 minutes + memory problems
 - The final sampling was done on a cluster with > 20 cores and 1TB of RAM
 - 6 chains, 300 + 300 samples, all chains had fixed initial points
- Convergence properties
 - \hat{R} -values were good (≤ 1.015)
 - Posterior chains had no trends
 - Divergences and maximum treedepths hit ≈ 0
 - Effective sample size estimates ESS_{bulk} and ESS_{tail} were ≥ 320

Results - Model checking

- Pooled and hierarchical models were compared using PSIS-LOO cross-validation
 - Results favoured the hierarchical model slightly
 - All pareto-k estimates were well under 0.5

	Pooled	Hierarchical
ELPD-LOO	7750	7761
p_{eff}	8.8	10.3

M8 posterior draws from the hierarchical model



colour

— M8 nut measured

— M8 Posterior draws

Results - Numerical values

- First nut-pendulum seems to have larger air-resistance
 - Uncertainty is quite large, differences might be due to imperfect oscillation

Pooled model		
Parameter	Mean	90%-quantile
Mass m [g]	13	[6, 21]
Air-resistance μ [$\frac{kg \times m}{s}$]	2.2×10^{-4}	$[1.1 \times 10^{-4}, 3.5 \times 10^{-4}]$
String length l [m]	0.171	[0.170, 0.172]
Gravity g [m/s^2]	9.81	[9.78, 9.86]
Hierarchical model		
Air-resistance M8 μ_1 [$\frac{kg \times m}{s}$]	2.4×10^{-4}	$[1.3 \times 10^{-4}, 3.7 \times 10^{-4}]$
Air-resistance M10 μ_2 [$\frac{kg \times m}{s}$]	2.0×10^{-4}	$[0.9 \times 10^{-4}, 3.3 \times 10^{-4}]$

Summary

- The aim of the project was to fit ODE to simple nut pendulum data using Stan
 - Ill-posed problem with 4 system parameters and 7 parameters in total
- Two models were used: Pooled and hierarchical
 - Models showed only minimal differences
 - Model comparison with LOO-CV supports this
 - Both models were computationally heavy
- The damping was different between the nuts
 - Though with so small difference that the root cause cannot be certainly assigned
 - Either real difference in the nut properties or difference between the two measurements



Additional material

Additional material - ODE

- System can be described with one coordinate θ
 - Equivalently with x-coordinate using small angle sine:


$$\sin(\theta) = \frac{x}{l} \Rightarrow \theta \approx \frac{x}{l}$$

- The ODE is then a 2nd order linear ODE

$$\ddot{x} = -\frac{\mu}{ml}\dot{x} - \frac{g}{l}x$$


- Assumptions:
 - θ is small, air-resistance is proportional to velocity

$$F_{\mu} = \mu\dot{x}$$



Additional material - Convergence

- There were no problems with convergence
 - Posterior chains had no trends



Additional material - Stan code

```
functions {  
  
  // Function that encapsulates the dynamics of the system  
  vector dx_dt(real t, vector x, real m, real l, real mu, real g) {  
    vector[2] dxdt;  
    dxdt[1] = x[2];  
    dxdt[2] = (-mu/(m*l))*x[2] - (g/l)*x[1];  
    return dxdt;  
  }  
  
  // Function that solves the ODE and returns the position array  
  real[] simulate_x(int N, real[] t, vector s0, real m,  
                    real l, real mu, real g){  
    vector[2] z[N-1] = ode_rk45(dx_dt, s0, 0.0, t[2:N], m, l, mu, g);  
    array[1] real x0;  
    x0[1] = s0[1];  
    real x_sim[N] = append_array(x0, z[,1]);  
    return x_sim;  
  }  
}
```


Additional material - Stan code

```
// The problem data  
data {  
  int<lower=1> N;  
  matrix[N,2] x;  
  array[N,2] real ts;  
  matrix[8, 2] pr_params; // Matrix for prior parameters  
}
```

Additional material - Stan code

```
// The problem parameters
parameters {
  vector[2] state01;
  vector[2] state02;
  real<lower=0> sigma;
  real<lower=0.000001> m; // Prevent division by zero issues
  real<lower=0.0001> l; // Prevent division by zero issues
  real<lower=0> mu;
  real<lower=0> g;
}
```

Additional material - Stan code

```
// The ODE-transformation from params to simulated data  
transformed parameters {  
  real x_sim1[N] = simulate_x(N, ts[,1], state01, m, l, mu, g);  
  real x_sim2[N] = simulate_x(N, ts[,2], state02, m, l, mu, g);  
}
```

Additional material - Stan code

```
// The actual model
model {
  // The priors
  m ~ normal(pr_params[1,1], pr_params[1,2]);
  l ~ normal(pr_params[2,1], pr_params[2,2]);
  mu ~ uniform(pr_params[3,1], pr_params[3,2]);
  g ~ normal(pr_params[4,1], pr_params[4,2]);
  state01[1] ~ normal(pr_params[5,1], pr_params[5,2]); // Position of nut 1
  state01[2] ~ normal(pr_params[7,1], pr_params[7,2]); // Velocity of nut 1
  state02[1] ~ normal(pr_params[6,1], pr_params[6,2]); // Position of nut 2
  state02[2] ~ normal(pr_params[7,1], pr_params[7,2]); // Velocity of nut 2
  sigma ~ normal(pr_params[8,1], pr_params[8,2]);

  // The likelihood model
  for (t in 1:N) {
    x[t,1] ~ normal(x_sim1[t], sigma);
    x[t,2] ~ normal(x_sim2[t], sigma);
  }
}
```

Additional material - Stan code

```
// Computing the needed quantities for LOO-CV and other statistics
generated quantities {
  // Posterior draws
  real x_draws1[N] = simulate_x(N, ts[,1], state01, m, l, mu, g);
  real x_draws2[N] = simulate_x(N, ts[,2], state02, m, l, mu, g);

  // Log-likelihoods
  matrix[N,2] log_lik;
  for (t in 1:N){
    log_lik[t,1] = normal_lpdf(x[t,1] | x_sim1[t], sigma);
    log_lik[t,2] = normal_lpdf(x[t,2] | x_sim2[t], sigma);
  }
}
```

